

PitchXPO Conclave 2026

Talk Proposals for Review

Submitted by **Uday Nagulavancha**

Ex-AWS Senior Technical Program Manager and FinOps Certified Practitioner who has helped Fortune 500 companies turn runaway AI spend into measurable value, delivering over \$150 million in savings and building AI cost-governance infrastructure that proves AI investments pay off.

Two 30-minute talk options are proposed below. Each includes a title, full abstract, and the high-level topics the session covers, for the committee's review and selection.

TOPIC 1 — PRIMARY SUBMISSION

Tokenomics: The New Unit Economics Every AI Founder and Investor Needs to Understand

Abstract

Every AI feature now carries a real, variable cost behind every single user interaction — and the math behind that cost is stranger than most people building AI products realize. AI spend is linear in tokens, never exponential, yet bills routinely blow up 30x when usage only grows 10x — because what actually drifts is tokens-per-user, not the pricing curve itself. That gap is also a margin problem: Bessemer's research shows AI-native companies running 50–60% gross margins versus 80–90% for classic SaaS, on the same measure. This talk introduces Tokenomics — the discipline of knowing what a query, a task, or a tenant actually costs before you ever set a customer price — through a practical three-stage framework: Inform (see the spend), Optimize (cut it — caching, batching, and the commitment math that actually pays off), and Operate (govern it, through a real story of an AI agent that cost 21x more than it should have, and the five guardrails that prevent it). Along the way: why Salesforce launched Agentforce on flat \$2-per-conversation pricing and had to add two more ways to pay within months, after customers couldn't predict their bills — and how Workday and ServiceNow landed on the same hybrid shape independently. Attendees leave with a live, interactive token-cost calculator to try themselves, a pricing checklist for founders, and — for investors — a new diligence question: ask what a startup's tokenomics and guardrails look like before you ask what they charge.

High-Level Topics Covered

- The token economics problem — why AI margins don't behave like classic SaaS margins

- The input vs. output cost asymmetry, and why it catches builders off guard
- The real reason AI bills explode: tokens-per-user drift, not exponential pricing
- A practical framework — Inform, Optimize, Operate — for managing AI unit economics
- A cautionary tale: how an ungoverned AI agent can cost 20x more than expected, and the five real guardrails that prevent it
- Live interactive demo: a token-cost calculator the audience uses on their own phones
- Takeaways: a pricing checklist for founders, and a new diligence question for investors

Format: 30-minute talk, including a live audience interactive segment.

TOPIC 2 — SECONDARY SUBMISSION

The ROI Illusion: Why Gartner Predicts 40% of Agentic AI Projects Get Cancelled — and How the Survivors Measure Value Differently

Abstract

Across enterprises today, AI infrastructure spend is compounding faster than anyone can attribute it to revenue — and Gartner predicts more than 40% of agentic AI projects will be cancelled by the end of 2027, not because the technology failed, but because nobody could prove the value. This talk breaks down why “what's our ROI on AI” is often the wrong question, and what actually separates AI investments that survive their first budget review from the 40% that don't. Drawing on Salesforce, Workday, and ServiceNow's own public earnings disclosures alongside original research on model economics, attendees leave with a concrete framework: reframing cost-per-prompt into cost-per-completed-task, sequencing visibility before optimization, and the specific signals — used by enterprise leaders and investors alike — that separate real AI value from AI theater.

High-Level Topics Covered

- Why Gartner predicts 40%+ of agentic AI projects will be cancelled by 2027
- Reframing the ROI question: cost-per-completed-task vs. cost-per-prompt
- The FinOps sequencing principle: visibility before optimization
- What public AI-revenue disclosures reveal — and hide — at Salesforce, Workday, and ServiceNow
- A three-question framework to separate real AI value from AI theater
- A worked example: one AI investment that survived, and one that got cut

Format: 30-minute talk, including a live audience interactive segment.